

Alexandria University

Faculty of Engineering

Computer and Systems Engineering Department

Assignment #1: Heart Failure Classification Problem

**CSE: Pattern Recognition Assigned: Sun, Mar. 02, 2025 Due: Tue, Mar.
11, 2025**

Submitted by:

Omar Hani Bishr - 21010891

Mohamed Hassan Mohamed Qabari Hassan - 21011115

Mohamed Mohamed Ibrahim Ali Hassan – 21011211

Instructor: Prof. Dr. Marwan Torki

TA: Eng. Ismail El-Yamany

Approach

Dataset Preparation and Preprocessing

We used the **ORL Face Dataset**, which contains grayscale images of 40 individuals, each with 10 images of size 92×112 pixels. The following steps were taken to prepare the data:

1. Image Flattening:

Each image was converted into a 1D vector of size $10304 = 92 \times 112$, and all 400 images were stacked to form a **data matrix** of shape **(400, 10304)**.

2. Label Vector:

Labels were extracted from folder names (e.g., s1, s2, ..., s40) and stored in a vector of shape **(400,)** with values from 1 to 40, representing subject IDs.

3. Normalization:

All pixel values were normalized to the range $[0, 1]$ by dividing by 255.

4. Train/Test Split (Per Subject):

To ensure equal subject representation:

- a. Odd-indexed images (1st, 3rd, 5th, 7th, 9th) were used for **training** — 5 images per subject.
- b. Even-indexed images (2nd, 4th, 6th, 8th, 10th) were used for **testing** — 5 images per subject.
- c. Final shapes:

- i. Training set: **(200, 10304)**
- ii. Test set: **(200, 10304)**
- iii. Labels: integers from 1 to 40

Principal Component Analysis (PCA)

To reduce dimensionality while preserving facial structure, we implemented PCA from scratch with the following steps:

1. PCA Class Implementation:

- a. `fit(X)`:
 - i. Centers the data and computes the **covariance matrix**.
 - ii. Extracts **eigenvalues and eigenvectors**, then sorts them in descending order.
 - iii. Saves them for later use.
- b. `reduce_dimensions(X, alpha)`:
 - i. Selects the **minimum number of components** such that the retained variance $\geq \alpha$ (e.g., $0.9 = 90\%$).
 - ii. Projects the data onto the selected eigenvectors (PCA space).
 - iii. Returns the reduced representation.
- c. `visualize_reconstruction(...)`:
 - i. Reconstructs and visualizes faces from reduced PCA space to compare with originals.

2. Training and Testing Transformation:

- a. Applied PCA to both **training** and **testing** sets separately for multiple alpha values: **[0.80, 0.85, 0.90, 0.95]**.
- b. Printed the number of principal components retained for each alpha.

KMeans Clustering on PCA-Reduced Data

To group similar facial features without supervision, we implemented a **custom KMeans algorithm** and evaluated clustering accuracy:

1. Custom KMeans Implementation:

- a. Randomly initialized K centroids from training data.
- b. Iteratively:
 - i. Assigned data points to the nearest centroid.
 - ii. Updated centroids based on the mean of assigned points.
 - iii. Checked for convergence using a tolerance threshold.

2. Cluster-to-Label Mapping:

- a. Mapped each cluster to the **most common true label** in that cluster using Counter.
- b. Predicted labels were then compared to true labels to compute clustering **accuracy**.

3. Experiment Setup:

- a. Ran KMeans on each reduced training set for PCA $\alpha \in \{0.80, 0.85, 0.90, 0.95\}$.
- b. Tested with cluster counts $K \in \{20, 40, 60\}$.
- c. Calculated accuracy using mapped labels.

4. Visualization:

- a. Created two line plots:
 - i. **Accuracy vs K** for each α .
 - ii. **Accuracy vs Alpha** for each K .
- b. Helped visualize the tradeoff between dimensionality reduction and clustering performance.

Gaussian Mixture Model (GMM)

Objective

To implement a robust and numerically stable Gaussian Mixture Model (GMM) from scratch that can cluster high-dimensional data using the Expectation-Maximization (EM) algorithm with multiple initializations and early stopping.

Key Components of the Approach

1. Initialization Strategy

- a. Means initialized using random data samples with added noise to avoid duplicates.

- b. Covariances initialized using scaled identity matrices based on data variance.
- c. Weights initialized with random non-uniform values to break symmetry.

2. EM Algorithm

- a. **E-step**: Calculates responsibilities using log-probabilities for numerical stability.
- b. **M-step**: Updates weights, means, and covariances using weighted statistics.
- c. Covariance matrices are regularized to prevent singularities and ensure positive definiteness.

3. Numerical Stability Techniques

- a. Uses `logsumexp` to prevent underflow in soft assignments.
- b. Adds `reg_covar` and additional regularization in exceptional cases (e.g., low determinant).
- c. Handles singular matrix errors via fallback mechanisms in both E-step and scoring.

4. Multiple Initializations & Model Selection

- a. Runs `n_init` times with different initializations.
- b. Selects the model with the highest log-likelihood score on the training data.

5. Convergence & Early Stopping

- a. Monitors log-likelihood difference.

- b. Stops early if no improvement for `early_stop` consecutive iterations.

6. Prediction & Scoring

- a. Predicts cluster labels using maximum posterior (responsibility).
- b. Computes total log-likelihood of data under the current GMM parameters.

Approach for Dimensionality Reduction with Autoencoder Embeddings

This **convolutional autoencoder (CAE)** compresses input images into a **low-dimensional embedding** ($7 \times 6 \times 1 = 42\text{D}$), then reconstructs them.

Key Steps:

1. Encoder (Dimensionality Reduction)

- Stacks **convolutional layers + max-pooling** to progressively downsample the image.
- Final embedding: **$7 \times 6 \times 1$ tensor** (42D flattened).

2. Decoder (Reconstruction)

- Uses **transposed convolutions** to upsample the embedding back to the original size.

3. Usage for Embeddings

- After training, discard the decoder.
- Use $\text{encoder}(x)$ to extract **compressed embeddings** (42D vectors) for downstream tasks (e.g., clustering, classification).

Why It Works

- The bottleneck layer ($7 \times 6 \times 1$) forces the network to learn a **compact representation**.
- Embeddings retain critical features while reducing dimensionality.

Results

Using PCA

PCA Reconstruction (alpha=0.8)

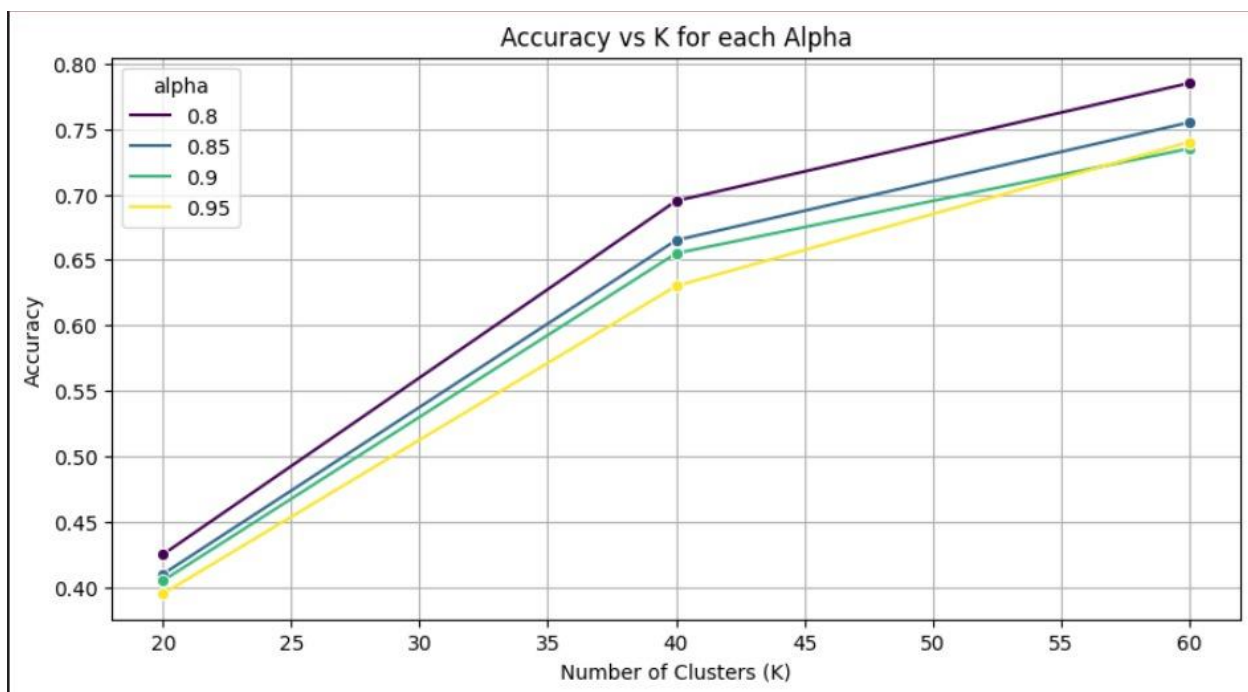
Original

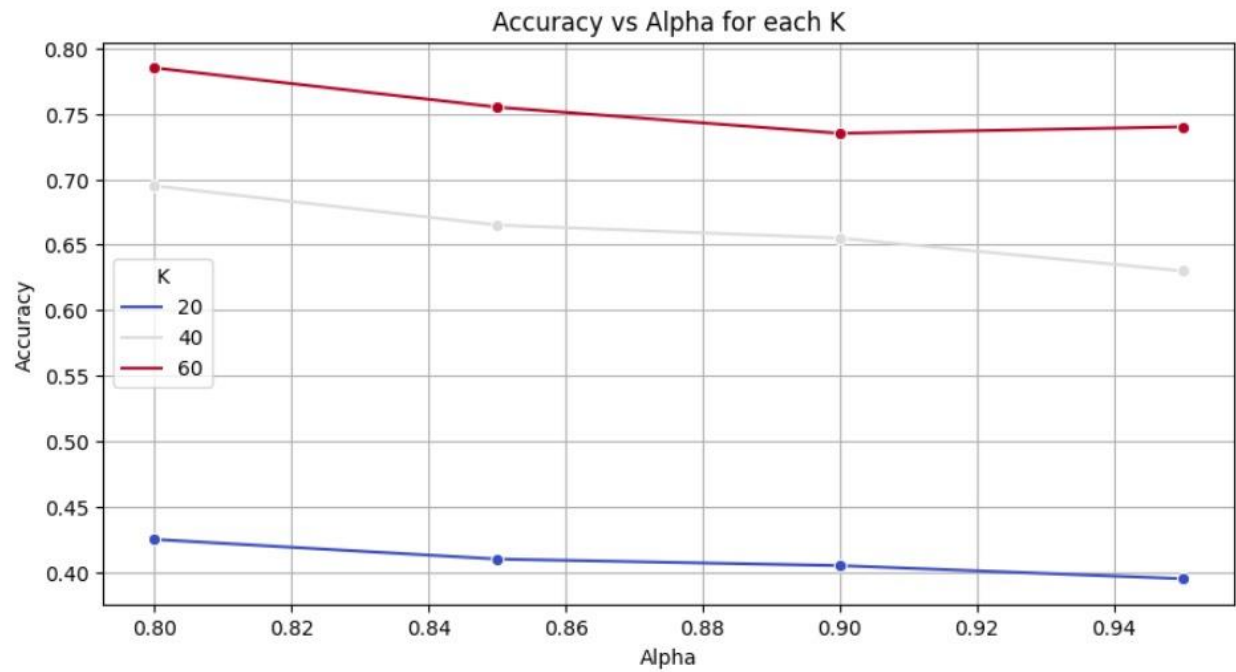


Reconstructed
MSE: 0.0064



K-means



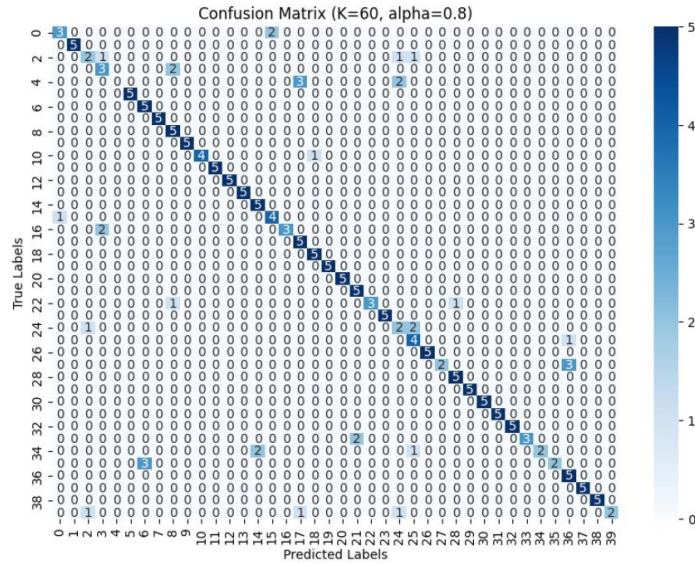


Best parameters -> Alpha: 0.8, K: 60

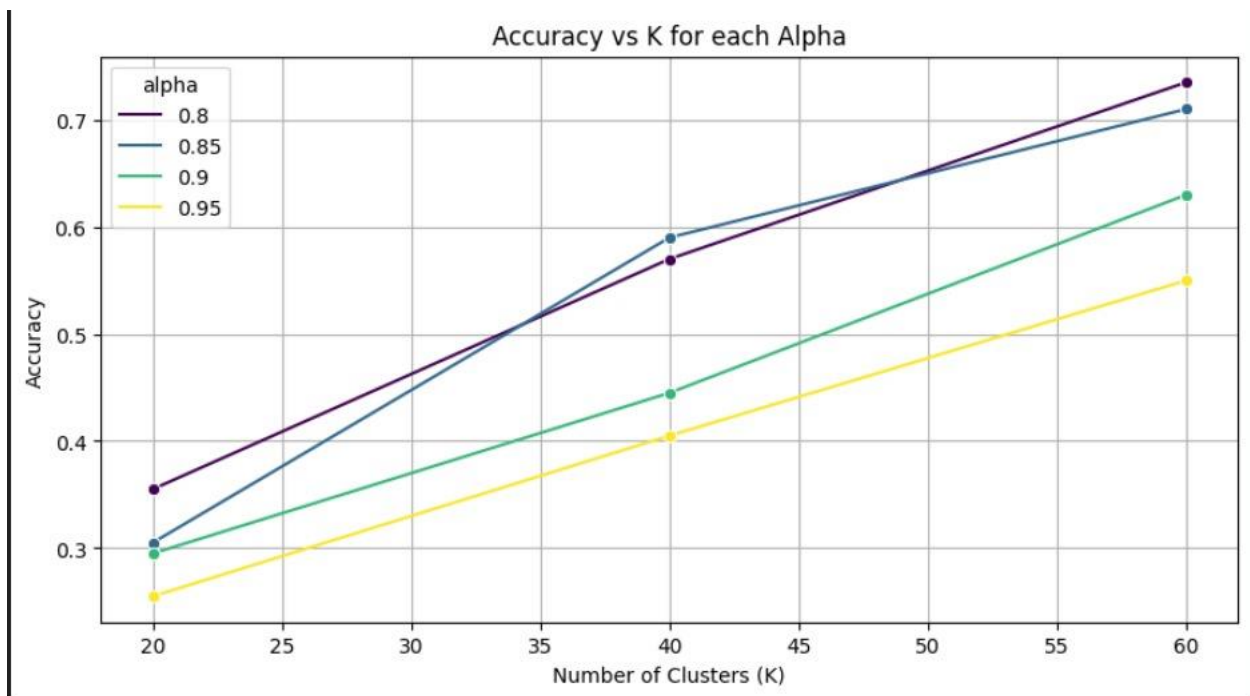
Results for K=60, alpha=0.8:

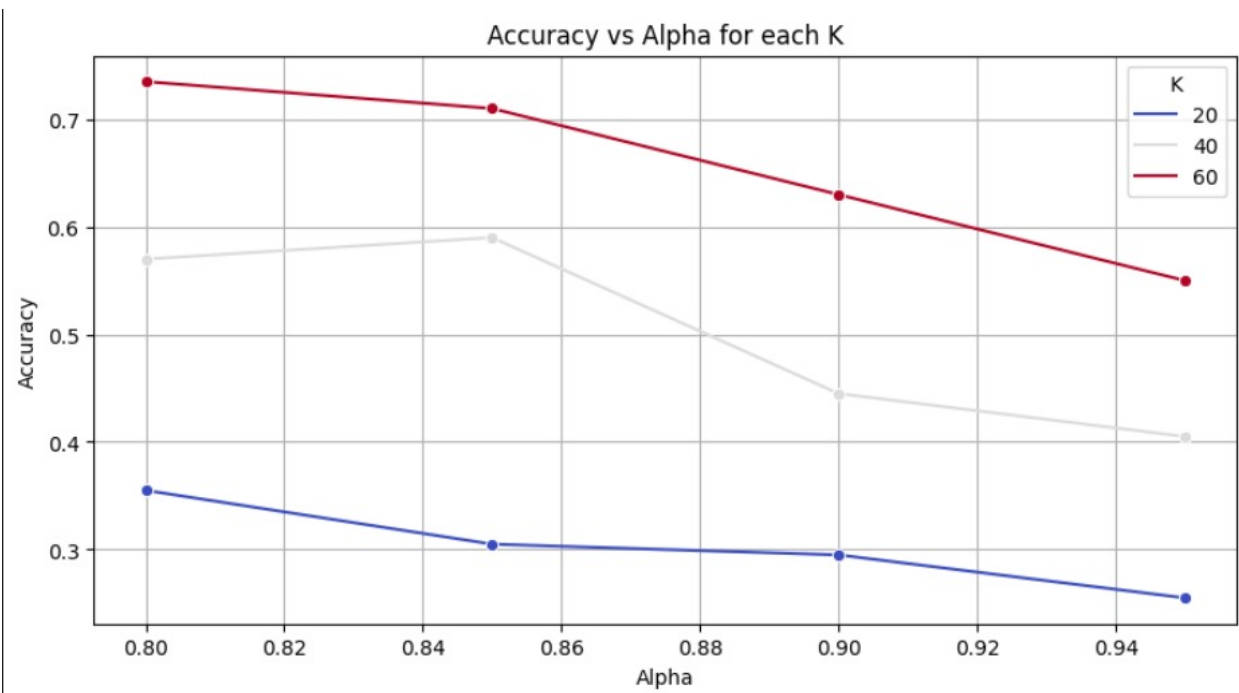
→ Accuracy: 0.8200

→ F1 Score: 0.8060



GMM






Results for $K=60$, $\alpha=0.8$:

→ Accuracy: 0.7650

→ F1 Score: 0.7524




PCA-based Clustering (K = 60, $\alpha = 0.8$)

Algorithm	Accuracy	F1 Score
KMeans	0.8200	0.8060
GMM	0.7650	0.7524

Using Auto-Encoder

Test Loss: 0.0073




K-means

Results for K=60:

→ Accuracy: 0.7700

→ F1 Score: 0.7452

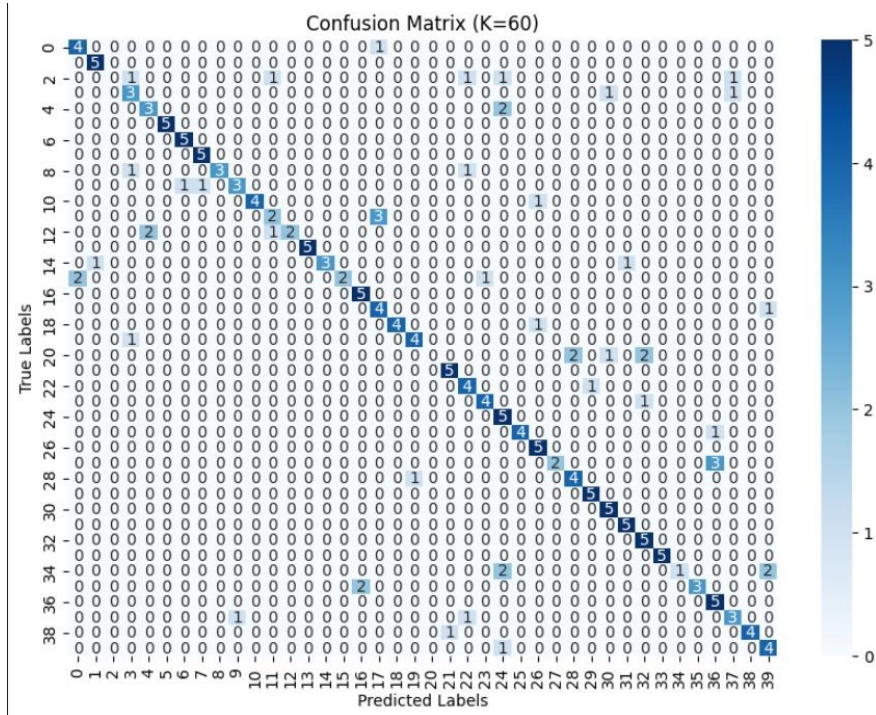


GMM

Results for K=60:

→ Accuracy: 0.7450

→ F1 Score: 0.7194



K: 20 → Accuracy: 0.4200

K: 40 → Accuracy: 0.6200

K: 60 → Accuracy: 0.7400

Autoencoder-based Clustering (K = 60)

Algorithm	Accuracy	F1 Score
KMeans	0.7700	0.7452
GMM	0.7450	0.7194

Analysis

Both in K-means & GMM

Can you find a relation between alpha and classification accuracy?
as alpha increases, accuracy decreases

Can you find a relation between K and classification accuracy?
as K increases, classification accuracy increases.

Using PCA

K-means

Best parameters -> Alpha: 0.8, K: 60

GMM

Best parameters -> Alpha: 0.8, K: 60

Using Auto-Endor

K-means

Best parameters -> K: 60

GMM

Best parameters -> K: 60

